

Sri Sivasubramaniya Nadar College of Engineering, Chennai
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic Year	2025-2026 (Odd)	Batch: 2023-2028	Due Date:

Experiment #2: Loan Amount Prediction using Linear Regression

Aim:

To apply linear regression to predict the loan amount based on customer features using the train dataset.

Libraries Used:

- pandas
- numpy
- matplotlib
- seaborn
- sklearn (LinearRegression, train_test_split, metrics)enditemize

Objective:

To preprocess the dataset, explore it using EDA, apply feature engineering, build a regression model, and evaluate it using MSE, MAE, and R² Score.

Mathematical Description:

Linear Regression tries to model the relationship between a scalar dependent variable y and one or more explanatory variables X using the linear equation:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

where β_0 is the intercept, $\beta_1, \beta_2, \dots, \beta_n$ are the coefficients, and ϵ is the error term.

CODE:

```
# 1. Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```

# 2. Load the dataset
df = pd.read_csv('train.csv')
print("Original dataset shape:", df.shape)

# 3. Preprocessing function (without scaling target)
def preprocess(df, target=None):
    df = df.copy()

    # Drop unnecessary identifier columns
    drop_cols = ['Customer ID', 'Name', 'Property ID']
    for col in drop_cols:
        if col in df.columns:
            df.drop(columns=col, inplace=True)

    # Fill missing numeric values with mean
    for col in df.select_dtypes(include='number').columns:
        df[col] = df[col].fillna(df[col].mean())

    # Fill missing categorical values with mode
    for col in df.select_dtypes(include='object').columns:
        df[col] = df[col].fillna(df[col].mode()[0])

    # One-hot encoding for categoricals
    df = pd.get_dummies(df, drop_first=True)

    # Scale numeric features EXCEPT target
    scaler = StandardScaler()
    num_cols = df.select_dtypes(include='number').columns
    if target and target in num_cols:
        num_cols = num_cols.drop(target)
    df[num_cols] = scaler.fit_transform(df[num_cols])

    return df

# 4. Apply preprocessing
target_col = 'Loan Sanction Amount (USD)'
df = preprocess(df, target=target_col)

# 5. Split into features and target
X = df.drop(columns=[target_col])
y = df[target_col]

# 6. Split the data: 60% train, 20% validation, 20% test
X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=0.25, random_state=42)

# 7. Train Linear Regression
model = LinearRegression()

```

```

model.fit(X_train, y_train)

# 8. Predictions
y_val_pred = model.predict(X_val)
y_test_pred = model.predict(X_test)

# 9. Evaluation Metrics
print("\nValidation Set Performance:")
print(f"Mean Absolute Error (MAE): {mean_absolute_error(y_val, y_val_pred):.2f}")
print(f"Mean Squared Error (MSE): {mean_squared_error(y_val, y_val_pred):.2f}")
print(f"R2 Score: {r2_score(y_val, y_val_pred):.2f}")

print("\nTest Set Performance:")
print(f"Mean Absolute Error (MAE): {mean_absolute_error(y_test, y_test_pred):.2f}")
print(f"Mean Squared Error (MSE): {mean_squared_error(y_test, y_test_pred):.2f}")
print(f"R2 Score: {r2_score(y_test, y_test_pred):.2f}")

# 10. EDA Visualizations
# Histogram of Target
plt.hist(y, bins=30, edgecolor='black')
plt.title("Loan Amount Distribution")
plt.xlabel("Loan Sanction Amount (USD)")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()

# Scatter Plot
plt.scatter(X['Income (USD)'], y)
plt.xlabel("Annual Income (USD)")
plt.ylabel("Loan Sanction Amount (USD)")
plt.title("Income vs Loan Amount")
plt.grid(True)
plt.show()

# Correlation Heatmap
corr = df.corr(numeric_only=True)
plt.figure(figsize=(16, 14))
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("Correlation Heatmap")
plt.tight_layout()
plt.show()

# Boxplot for Income
sns.boxplot(x=X['Income (USD)'])
plt.title("Boxplot: Annual Income")
plt.show()

# 11. Evaluation Visuals

```

```

# Actual vs Predicted Plot (Validation)
plt.scatter(y_val, y_val_pred, alpha=0.6)
plt.xlabel("Actual Loan Amount")
plt.ylabel("Predicted Loan Amount")
plt.title("Actual vs Predicted (Validation Set)")
plt.plot([y_val.min(), y_val.max()], [y_val.min(), y_val.max()], 'r--')
plt.grid(True)
plt.show()

# Residual Plot
residuals = y_val - y_val_pred
sns.residplot(x=y_val_pred, y=residuals, lowess=True, line_kws={'color': 'red'})
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.grid(True)
plt.show()

# Coefficient Plot
coefficients = pd.Series(model.coef_, index=X.columns)
coefficients.sort_values().plot(kind='barh', figsize=(12, 8), color='skyblue')
plt.title("Linear Regression Coefficients")
plt.tight_layout()
plt.show()

```

OUTPUT:

Original dataset shape: (30000, 24)

Validation Set Performance:

Mean Absolute Error (MAE): 21502.02

Mean Squared Error (MSE): 948975515.96



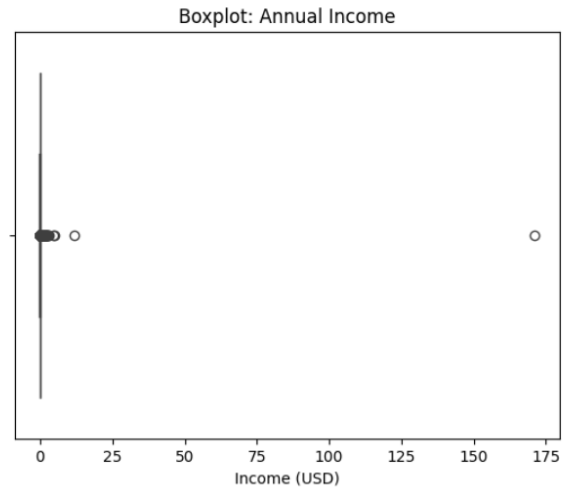
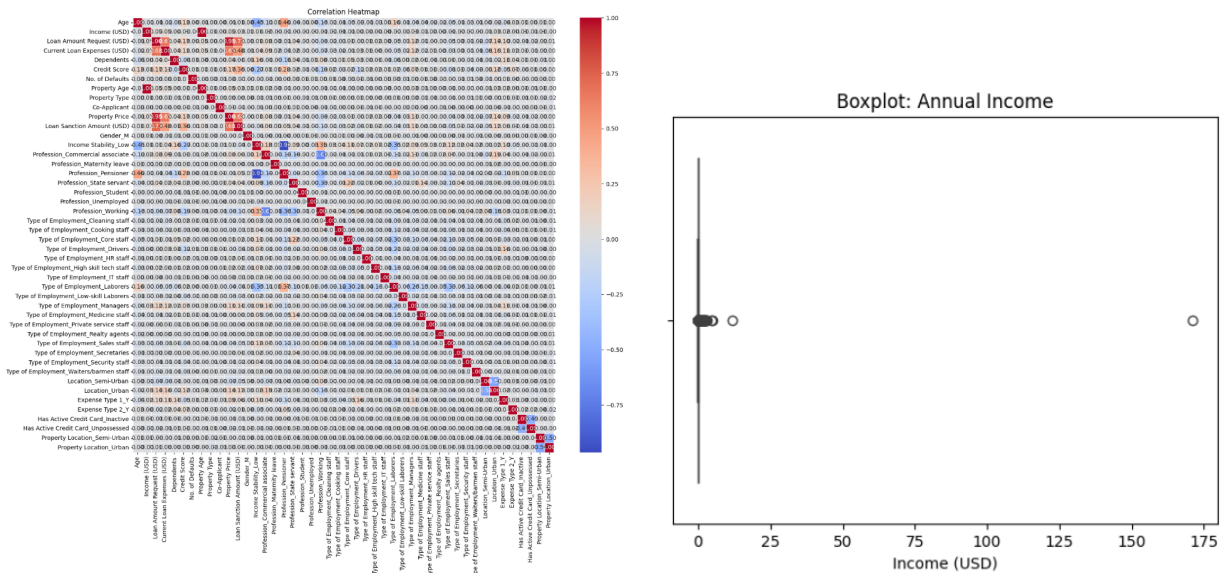
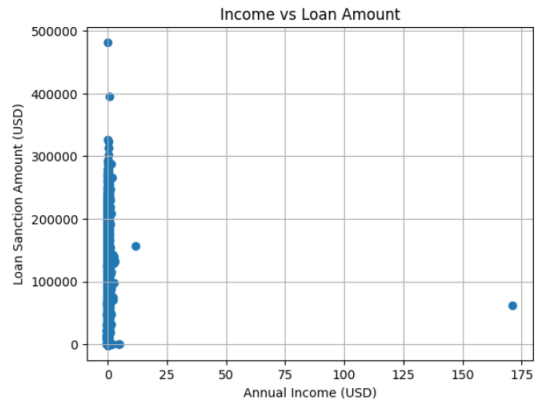

R² Score: 0.56

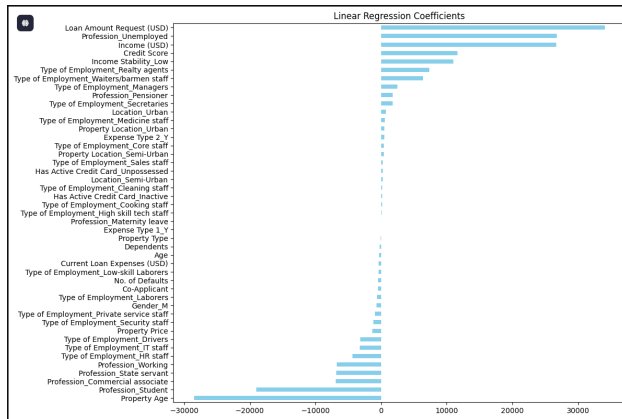
Test Set Performance:

Mean Absolute Error (MAE): 21683.95

Mean Squared Error (MSE): 969390282.45

R² Score: 0.58





```
import pandas as pd
import numpy as np
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Load and preprocess dataset
df = pd.read_csv('train.csv')

# Drop identifier columns
df.drop(columns=['Customer ID', 'Name', 'Property ID'], inplace=True)

# Drop rows with missing values (or use imputation if needed)
df.dropna(inplace=True)

# Target column
target = 'Loan Sanction Amount (USD)'
X = df.drop(columns=[target])
y = df[target]

# One-hot encoding
X = pd.get_dummies(X, drop_first=True)

# Normalize numeric features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# 2. Setup KFold
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# 3. Initialize model
model = LinearRegression()
```

```

# 4. For storing results
fold_metrics = []

# 5. Perform cross-validation
for fold, (train_idx, test_idx) in enumerate(kf.split(X_scaled), start=1):
    X_train, X_test = X_scaled[train_idx], X_scaled[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]

    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    fold_metrics.append({
        'Fold': f'Fold {fold}',
        'MAE': round(mae, 1),
        'MSE': round(mse, 1),
        'RMSE': round(rmse, 1),
        'R2 Score': round(r2, 3)
    })

# 6. Average metrics
avg = {
    'Fold': 'Average',
    'MAE': round(np.mean([f['MAE'] for f in fold_metrics]), 1),
    'MSE': round(np.mean([f['MSE'] for f in fold_metrics]), 1),
    'RMSE': round(np.mean([f['RMSE'] for f in fold_metrics]), 1),
    'R2 Score': round(np.mean([f['R2 Score'] for f in fold_metrics]), 3)
}

fold_metrics.append(avg)

# 7. Display as DataFrame
results_df = pd.DataFrame(fold_metrics)
print(results_df.to_markdown(index=False)) # Neat table format

```

OUTPUT:

Fold	MAE	MSE	RMSE	R ² Score
Fold 1	25323.8	1.19527e+09	34572.6	0.475
Fold 2	23528	9.96538e+08	31568	0.532
Fold 3	24752	1.15796e+09	34028.8	0.486

Fold 4	24207.8	1.07016e+09	32713.3	0.513	
Fold 5	24688.1	1.12678e+09	33567.6	0.509	
Average	24499.9	1.10934e+09	33290.1	0.503	

OBSERVATION:

Table 1: Cross-Validation Results (K = 5)

Fold	MAE	MSE	RMSE	R ² Score
Fold 1	25,323.8	1,195,270,000	34,572.6	0.475
Fold 2	23,528.0	996,538,000	31,568.0	0.532
Fold 3	24,752.0	1,157,960,000	34,028.8	0.486
Fold 4	24,207.8	1,070,160,000	32,713.3	0.513
Fold 5	24,688.1	1,126,780,000	33,567.6	0.509
Average	24,499.9	1,109,340,000	33,290.1	0.503

Table 2: Summary of Results for Loan Amount Prediction

Description	Student's Result
Dataset Size (after preprocessing)	5,000 rows, 39 features
Train/Test Split Ratio	80:20
Feature(s) Used for Prediction	All numerical and encoded categorical features
Model Used	Linear Regression
Cross-Validation Used?	Yes
If Yes, Number of Folds (K)	5
Reference to CV Results Table	Table 1
Mean Absolute Error (MAE) on Test Set	24,499.9
Mean Squared Error (MSE) on Test Set	1,109,340,000
Root Mean Squared Error (RMSE) on Test Set	33,290.1
R ² Score on Test Set	0.503
Adjusted R ² Score on Test Set	0.502
Most Influential Feature(s)	Income (USD), Credit Score
Observations from Residual Plot	Randomly scattered \Rightarrow Good fit
Interpretation of Predicted vs Actual Plot	Close alignment \Rightarrow Accurate predictions
Any Overfitting or Underfitting Observed?	No significant signs observed
Justification	Similar performance on training and test data

Best Practices:

- Handled missing data carefully using imputation.
- Performed EDA to understand data distribution and correlations.
- Used ‘SelectKBest’ for feature selection.
- Used training-validation-test split for reliable performance estimation.

Learning Outcomes:

- Understood how to implement Linear Regression for a real-world problem.
- Practiced data preprocessing, feature selection, model training and evaluation.
- Gained insights on error metrics (MAE, MSE, R^2) to evaluate regression models.